

Configure health probes for pods.

Overview

Implementing [health probes](#) in your application is a great way for Kubernetes to automate some tasks to improve availability in the event of an error.

Because OSM reconfigures application Pods to redirect all incoming and outgoing network traffic through the proxy sidecar, [httpGet](#) and [tcpSocket](#) health probes invoked by the kubelet will fail due to the lack of any mTLS context required by the proxy.

For health probes to continue to work as expected from within the mesh, OSM adds configuration to expose the probe endpoint via the proxy and rewrites the probe definitions for new Pods to refer to the proxy-exposed endpoint. All of the functionality of the original probe is still used, OSM simply fronts it with the proxy so the kubelet can communicate with it.

For HTTP probes, the following table shows the modified path and port for each probe type.

Probe	Path
Liveness	/osm-liveness-probe
Readiness	/osm-readiness-probe
Startup	/osm-startup-probe

HTTPS and [tcpSocket](#) probes will have their ports modified the same way as HTTP probes. For HTTPS probes, the path is left unchanged.

Only predefined [httpGet](#) and [tcpSocket](#) probes are modified. If a probe is undefined, one will not be added in its place. `exec` probes (including those using `grpc_health_probe`) are never modified and will continue to function as expected as long as the command does not require network access outside of [localhost](#).

Examples

The following examples show how OSM handles health probes for Pods in a mesh.

HTTP

Consider a Pod spec defining a container with the following [livenessProbe](#):

```
livenessProbe:
  httpGet:
    path: /liveness
    port: 14001
    scheme: HTTP
```

When the Pod is created, OSM will modify the probe to be the following:

```
livenessProbe:
```

```
httpGet:
  path: /osm-liveness-probe
  port: 15901
  scheme: HTTP
```

The Pod's proxy will contain the following Envoy configuration.

An Envoy cluster which maps to the original probe port 14001:

```
{
  "cluster": {
    "@type": "type.googleapis.com/envoy.config.cluster.v3.Cluster",
    "name": "liveness_cluster",
    "type": "STATIC",
    "connect_timeout": "1s",
    "load_assignment": {
      "cluster_name": "liveness_cluster",
      "endpoints": [
        {
          "lb_endpoints": [
            {
              "endpoint": {
                "address": {
                  "socket_address": {
                    "address": "0.0.0.0",
                    "port_value": 14001
                  }
                }
              }
            ]
          }
        ]
      }
    },
    "last_updated": "2021-03-29T21:02:59.086Z"
  }
}
```

A listener for the new proxy-exposed HTTP endpoint at `/osm-liveness-probe` on port 15901 mapping to the cluster above:

```
{
  "listener": {
    "@type": "type.googleapis.com/envoy.config.listener.v3.Listener",
    "name": "liveness_listener",
    "address": {
      "socket_address": {
        "address": "0.0.0.0",
```

```

    "port_value": 15901
  },
  "filter_chains": [
    {
      "filters": [
        {
          "name": "envoy.filters.network.http_connection_manager",
          "typed_config": {
            "@type":
"type.googleapis.com/envoy.extensions.filters.network.http_connection_manager.v3.HttpConnec
tionManager",
            "stat_prefix": "health_probes_http",
            "route_config": {
              "name": "local_route",
              "virtual_hosts": [
                {
                  "name": "local_service",
                  "domains": [
                    "*"
                  ],
                  "routes": [
                    {
                      "match": {
                        "prefix": "/osm-liveness-probe"
                      },
                      "route": {
                        "cluster": "liveness_cluster",
                        "prefix_rewrite": "/liveness"
                      }
                    }
                  ]
                }
              ]
            },
            "http_filters": [...],
            "access_log": [...]
          }
        }
      ]
    }
  ],
  "last_updated": "2021-03-29T21:02:59.092Z"
}

```

HTTPS

Consider a Pod spec defining a container with the following `livenessProbe`:

```
livenessProbe:
  httpGet:
    path: /liveness
    port: 14001
    scheme: HTTPS
```

When the Pod is created, OSM will modify the probe to be the following:

```
livenessProbe:
  httpGet:
    path: /liveness
    port: 15901
    scheme: HTTPS
```

The Pod's proxy will contain the following Envoy configuration.

An Envoy cluster which maps to the original probe port 14001:

```
{
  "cluster": {
    "@type": "type.googleapis.com/envoy.config.cluster.v3.Cluster",
    "name": "liveness_cluster",
    "type": "STATIC",
    "connect_timeout": "1s",
    "load_assignment": {
      "cluster_name": "liveness_cluster",
      "endpoints": [
        {
          "lb_endpoints": [
            {
              "endpoint": {
                "address": {
                  "socket_address": {
                    "address": "0.0.0.0",
                    "port_value": 14001
                  }
                }
              }
            }
          ]
        }
      ]
    }
  },
  "last_updated": "2021-03-29T21:02:59.086Z"
}
```

A listener for the new proxy-exposed TCP endpoint on port 15901 mapping to the cluster above:

```
{
  "listener": {
    "@type": "type.googleapis.com/envoy.config.listener.v3.Listener",
    "name": "liveness_listener",
    "address": {
      "socket_address": {
        "address": "0.0.0.0",
        "port_value": 15901
      }
    },
    "filter_chains": [
      {
        "filters": [
          {
            "name": "envoy.filters.network.tcp_proxy",
            "typed_config": {
              "@type":
"type.googleapis.com/envoy.extensions.filters.network.tcp_proxy.v3.TcpProxy",
              "stat_prefix": "health_probes",
              "cluster": "liveness_cluster",
              "access_log": [...]
            }
          }
        ]
      }
    ],
    "last_updated": "2021-04-07T15:09:22.704Z"
  }
}
```

tcpSocket

Consider a Pod spec defining a container with the following `livenessProbe`:

`livenessProbe:`

`tcpSocket:`

`port: 14001`

When the Pod is created, OSM will modify the probe to be the following:

`livenessProbe:`

`tcpSocket:`

`port: 15901`

The Pod's proxy will contain the following Envoy configuration.

An Envoy cluster which maps to the original probe port 14001:

```

{
  "cluster": {
    "@type": "type.googleapis.com/envoy.config.cluster.v3.Cluster",
    "name": "liveness_cluster",
    "type": "STATIC",
    "connect_timeout": "1s",
    "load_assignment": {
      "cluster_name": "liveness_cluster",
      "endpoints": [
        {
          "lb_endpoints": [
            {
              "endpoint": {
                "address": {
                  "socket_address": {
                    "address": "0.0.0.0",
                    "port_value": 14001
                  }
                }
              }
            }
          ]
        }
      ]
    }
  },
  "last_updated": "2021-03-29T21:02:59.086Z"
}

```

A listener for the new proxy-exposed TCP endpoint on port 15901 mapping to the cluster above:

```

{
  "listener": {
    "@type": "type.googleapis.com/envoy.config.listener.v3.Listener",
    "name": "liveness_listener",
    "address": {
      "socket_address": {
        "address": "0.0.0.0",
        "port_value": 15901
      }
    }
  },
  "filter_chains": [
    {
      "filters": [
        {
          "name": "envoy.filters.network.tcp_proxy",
          "typed_config": {

```

```

    "@type":
"type.googleapis.com/envoy.extensions.filters.network.tcp_proxy.v3.TcpProxy",
    "stat_prefix": "health_probes",
    "cluster": "liveness_cluster",
    "access_log": [...]
  }
}
]
}
],
},
"last_updated": "2021-04-07T15:09:22.704Z"
}

```

How to Verify Health of Pods in the Mesh

Kubernetes will automatically poll the health endpoints of Pods configured with startup, liveness, and readiness probes.

When a startup probe fails, Kubernetes will generate an Event (visible by `kubectl describe pod <pod name>`) and restart the Pod. The `kubectl describe` output may look like this:

```

...
Events:
  Type    Reason      Age    From          Message
  ----    -
Normal   Scheduled   17s    default-scheduler   Successfully assigned bookstore/bookstore-
v1-699c79b9dc-5g8zn to osm-control-plane
Normal   Pulled      16s    kubelet        Successfully pulled image
"openservicemesh/init:v0.10.0" in 26.5835ms
Normal   Created     16s    kubelet        Created container osm-init
Normal   Started     16s    kubelet        Started container osm-init
Normal   Pulling     16s    kubelet        Pulling image "openservicemesh/init:v0.10.0"
Normal   Pulling     15s    kubelet        Pulling image "envoyproxy/envoy-
alpine:v1.17.2"
Normal   Pulling     15s    kubelet        Pulling image
"openservicemesh/bookstore:v0.10.0"
Normal   Pulled      15s    kubelet        Successfully pulled image
"openservicemesh/bookstore:v0.10.0" in 319.9863ms
Normal   Started     15s    kubelet        Started container bookstore-v1
Normal   Created     15s    kubelet        Created container bookstore-v1
Normal   Pulled      14s    kubelet        Successfully pulled image "envoyproxy/envoy-
alpine:v1.17.2" in 755.2666ms
Normal   Created     14s    kubelet        Created container envoy
Normal   Started     14s    kubelet        Started container envoy

```

```
Warning Unhealthy 13s          kubelet          Startup probe failed: Get
"http://10.244.0.23:15903/osm-startup-probe": dial tcp 10.244.0.23:15903: connect: connection
refused
```

```
Warning Unhealthy 3s (x2 over 8s) kubelet          Startup probe failed: HTTP probe failed
with statuscode: 503
```

When a liveness probe fails, Kubernetes will generate an Event (visible by `kubectl describe pod <pod name>`) and restart the Pod. The `kubectl describe` output may look like this:

```
...
Events:
  Type    Reason      Age          From          Message
  ----    -
Normal    Scheduled   59s          default-scheduler Successfully assigned
bookstore/bookstore-v1-746977967c-jqjt4 to osm-control-plane
Normal    Pulling     58s          kubelet        Pulling image "openservicemesh/init:v0.10.0"
Normal    Created     58s          kubelet        Created container osm-init
Normal    Started     58s          kubelet        Started container osm-init
Normal    Pulled      58s          kubelet        Successfully pulled image
"openservicemesh/init:v0.10.0" in 23.415ms
Normal    Pulled      57s          kubelet        Successfully pulled image "envoyproxy/envoy-
alpine:v1.17.2" in 678.1391ms
Normal    Pulled      57s          kubelet        Successfully pulled image
"openservicemesh/bookstore:v0.10.0" in 230.3681ms
Normal    Created     57s          kubelet        Created container envoy
Normal    Pulling     57s          kubelet        Pulling image "envoyproxy/envoy-
alpine:v1.17.2"
Normal    Started     56s          kubelet        Started container envoy
Normal    Pulled      44s          kubelet        Successfully pulled image
"openservicemesh/bookstore:v0.10.0" in 20.6731ms
Normal    Created     44s (x2 over 57s) kubelet        Created container bookstore-v1
Normal    Started     43s (x2 over 57s) kubelet        Started container bookstore-v1
Normal    Pulling     32s (x3 over 58s) kubelet        Pulling image
"openservicemesh/bookstore:v0.10.0"
Warning Unhealthy 32s (x6 over 50s) kubelet        Liveness probe failed: HTTP probe
failed with statuscode: 503
Normal    Killing     32s (x2 over 44s) kubelet        Container bookstore-v1 failed liveness
probe, will be restarted
```

When a readiness probe fails, Kubernetes will generate an Event (visible with `kubectl describe pod <pod name>`) and ensure no traffic destined for Services the Pod may be backing is routed to the unhealthy Pod. The `kubectl describe` output for a Pod with a failing readiness probe may look like this:

```
...
Events:
  Type    Reason      Age          From          Message
  ----    -
```



```

Normal Scheduled 32s          default-scheduler Successfully assigned
bookstore/bookstore-v1-5848999cb6-hp6qg to osm-control-plane
Normal Pulling 31s           kubelet           Pulling image "openservicemesh/init:v0.10.0"
Normal Pulled 31s           kubelet           Successfully pulled image
"openservicemesh/init:v0.10.0" in 19.8726ms
Normal Created 31s          kubelet           Created container osm-init
Normal Started 31s          kubelet           Started container osm-init
Normal Created 30s          kubelet           Created container bookstore-v1
Normal Pulled 30s           kubelet           Successfully pulled image
"openservicemesh/bookstore:v0.10.0" in 314.3628ms
Normal Pulling 30s          kubelet           Pulling image
"openservicemesh/bookstore:v0.10.0"
Normal Started 30s          kubelet           Started container bookstore-v1
Normal Pulling 30s          kubelet           Pulling image "envoyproxy/envoy-
alpine:v1.17.2"
Normal Pulled 29s           kubelet           Successfully pulled image "envoyproxy/envoy-
alpine:v1.17.2" in 739.3931ms
Normal Created 29s          kubelet           Created container envoy
Normal Started 29s          kubelet           Started container envoy
Warning Unhealthy 0s (x3 over 20s) kubelet           Readiness probe failed: HTTP probe
failed with statuscode: 503

```

The Pod's `status` will also indicate that it is not ready which is shown in its `kubectl get pod` output. For example:

```

NAME                                READY STATUS RESTARTS AGE
bookstore-v1-5848999cb6-hp6qg 1/2 Running 0 85s

```

The Pods' health probes may also be invoked manually by forwarding the Pod's necessary port and using `curl` or any other HTTP client to issue requests. For example, to verify the liveness probe for the bookstore-v1 demo Pod, forward port 15901:

```
kubectl port-forward -n bookstore deployment/bookstore-v1 15901
```

Then, in a separate terminal instance, `curl` may be used to check the endpoint. The following example shows a healthy bookstore-v1:

```

$ curl -i localhost:15901/osm-liveness-probe
HTTP/1.1 200 OK
date: Wed, 31 Mar 2021 16:00:01 GMT
content-length: 1396
content-type: text/html; charset=utf-8
x-envoy-upstream-service-time: 1
server: envoy

```

```

<!doctype html>
<html itemscope="" itemtype="http://schema.org/WebPage" lang="en">
...
</html>

```

Known issues

- [#2207](#)

Troubleshooting

If any health probes are consistently failing, perform the following steps to identify the root cause:

1. Verify `httpGet` and `tcpSocket` probes on Pods in the mesh have been modified.

Startup, liveness, and readiness `httpGet` and `tcpSocket` probes must be modified by OSM in order to continue to function while in a mesh. Ports must be modified to 15901, 15902, and 15903 for liveness, readiness, and startup probes, respectively. Only HTTP (not HTTPS) probes will have paths modified in addition to be `/osm-liveness-probe`, `/osm-readiness-probe`, or `/osm-startup-probe`.

Also, verify the Pod's Envoy configuration contains a listener for the modified endpoint.

See the [examples above](#) for more details.

2. Determine if Kubernetes encountered any other errors while scheduling or starting the Pod.

Look for any errors that may have recently occurred with `kubect describe` of the unhealthy Pod. Resolve any errors and verify the Pod's health again.

3. Determine if the Pod encountered a runtime error.

Look for any errors that may have occurred after the container started by inspecting its logs with `kubect logs`. Resolve any errors and verify the Pod's health again.